

# Clase02Estadistica

November 18, 2024

Librerias

```
[1]: import os
os.chdir('C://Users//jaag2//OneDrive//PuraVidaAnalytics//00 - Clientes//03 - Texas Tech//07 - Cursos//01 - No Supervisado//03 - Estadistica')
```

```
[2]: import pandas as pd
import numpy as np
```

## 1 Leer Datos

```
[3]: carsdf = pd.read_csv("cars.csv")
carsdf
```

```
[3]:
```

	manufacturer_name	model_name	transmission	color	odometer_value	\
0	Subaru	Outback	automatic	silver	190000	
1	Subaru	Outback	automatic	blue	290000	
2	Subaru	Forester	automatic	red	402000	
3	Subaru	Impreza	mechanical	blue	10000	
4	Subaru	Legacy	automatic	black	280000	
...	...	...	...	...	...	
38526	Chrysler	300	automatic	silver	290000	
38527	Chrysler	PT Cruiser	mechanical	blue	321000	
38528	Chrysler	300	automatic	blue	777957	
38529	Chrysler	PT Cruiser	mechanical	black	20000	
38530	Chrysler	Voyager	automatic	silver	297729	

	year_produced	engine_fuel	engine_has_gas	engine_type	engine_capacity	\
0	2010	gasoline	False	gasoline	2.5	
1	2002	gasoline	False	gasoline	3.0	
2	2001	gasoline	False	gasoline	2.5	
3	1999	gasoline	False	gasoline	3.0	
4	2001	gasoline	False	gasoline	2.5	
...	...	...	...	...	...	
38526	2000	gasoline	False	gasoline	3.5	
38527	2004	diesel	False	diesel	2.2	
38528	2000	gasoline	False	gasoline	3.5	

38529	2001	gasoline	False	gasoline	2.0
38530	2000	gasoline	False	gasoline	2.4

	...	feature_1	feature_2	feature_3	feature_4	feature_5	feature_6	\
0	...	True	True	True	False	True	False	
1	...	True	False	False	True	True	False	
2	...	True	False	False	False	False	False	
3	...	False	False	False	False	False	False	
4	...	True	False	True	True	False	False	
...	...	...	...	...	...	...	...	
38526	...	True	False	False	True	True	False	
38527	...	True	False	False	True	True	False	
38528	...	True	False	False	True	True	False	
38529	...	True	False	False	False	False	False	
38530	...	False	False	False	False	False	False	

	feature_7	feature_8	feature_9	duration_listed
0	True	True	True	16
1	False	False	True	83
2	False	True	True	151
3	False	False	False	86
4	False	False	True	7
...	...	...	...	...
38526	False	True	True	301
38527	False	True	True	317
38528	False	True	True	369
38529	False	False	True	490
38530	False	False	True	632

[38531 rows x 30 columns]

```
[4]: carsdf.dtypes # Tipos de Datos
```

```
[4]: manufacturer_name    object
model_name                object
transmission              object
color                    object
odometer_value            int64
year_produced             int64
engine_fuel               object
engine_has_gas            bool
engine_type               object
engine_capacity           float64
body_type                 object
has_warranty              bool
state                    object
drivetrain                object
```

```

price_usd          float64
is_exchangeable    bool
location_region    object
number_of_photos   int64
up_counter         int64
feature_0          bool
feature_1          bool
feature_2          bool
feature_3          bool
feature_4          bool
feature_5          bool
feature_6          bool
feature_7          bool
feature_8          bool
feature_9          bool
duration_listed    int64
dtype: object

```

Generar un conjunto de estadísticos descriptivos del dataset:

```
[5]: carsdf.describe()
```

```

[5]:      odometer_value  year_produced  engine_capacity  price_usd \
count      38531.000000    38531.000000      38521.000000  38531.000000
mean      248864.638447      2002.943734         2.055161   6639.971021
std       136072.376530         8.065731         0.671178   6428.152018
min           0.000000      1942.000000         0.200000     1.000000
25%       158000.000000      1998.000000         1.600000   2100.000000
50%       250000.000000      2003.000000         2.000000   4800.000000
75%       325000.000000      2009.000000         2.300000   8990.000000
max      1000000.000000      2019.000000         8.000000  50000.000000

      number_of_photos  up_counter  duration_listed
count      38531.000000  38531.000000      38531.000000
mean           9.649062    16.306091        80.577249
std           6.093217    43.286933       112.826569
min            1.000000     1.000000         0.000000
25%            5.000000     2.000000        23.000000
50%            8.000000     5.000000        59.000000
75%           12.000000    16.000000        91.000000
max           86.000000   1861.000000       2232.000000

```

## 2 Medidas de Tendencia Central

```
[6]: carsdf.price_usd.mean()
```

```
[6]: 6639.971021255613
```

```
[7]: carsdf['price_usd'].mean()
```

```
[7]: 6639.971021255613
```

```
[8]: carsdf['price_usd'].quantile(0.25)
```

```
[8]: 2100.0
```

```
[9]: carsdf['price_usd'].std()
```

```
[9]: 6428.152018202915
```

```
[10]: carsdf['price_usd'].quantile(.50) #Mediana
```

```
[10]: 4800.0
```

```
[11]: carsdf['price_usd'].min()
```

```
[11]: 1.0
```

```
[12]: carsdf['price_usd'].max()
```


```
[12]: 50000.0
```

```
[13]: carsdf['price_usd'].mode()
```

```
[13]: 0    1500.0  
      Name: price_usd, dtype: float64
```

```
[14]: carsdf['price_usd'].plot.hist(bins=15)
```

```
[14]: <Axes: ylabel='Frequency'>
```




```
[15]: #Rango = valor max - valor min
rang = carsdf['price_usd'].max() - carsdf['price_usd'].min()
rang
```

```
[15]: 49999.0
```

```
[16]: import seaborn as sns
sns.boxplot(
    y      = 'price_usd',
    data   = carsdf
)
```

```
[16]: <Axes: ylabel='price_usd'>
```



```
[17]: #Quartiles
median = carsdf['price_usd'].median()
Q1 = carsdf['price_usd'].quantile(q=0.25) #toma el primer 25% de todos los datos
Q3 = carsdf['price_usd'].quantile(q=0.75)
min_val = carsdf['price_usd'].quantile(q=0)
max_val = carsdf['price_usd'].quantile(q=1)
print(min_val, Q1, median, Q3, max_val)
```

```
1.0 2100.0 4800.0 8990.0 50000.0
```

```
[18]: iqr = Q3 - Q1
iqr
```

```
[18]: 6890.0
```

## 2.1 Límites para detección de outliers (con datos simétricamente distribuidos)

```
[19]: minlimit = Q1 - 1.5*iqr
maxlimit = Q3 + 1.5*iqr
print(minlimit, maxlimit)
```

```
-8235.0 19325.0
```


```
[20]: minlimit
```

```
[20]: -8235.0
```

```
[21]: sns.histplot(carsdf['price_usd'])
```

```
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed in a  
future version. Convert inf values to NaN before operating instead.  
with pd.option_context('mode.use_inf_as_na', True):
```


```
[21]: <Axes: xlabel='price_usd', ylabel='Count'>
```



**Nota 1:** El valor es negativo porque se está aplicando una ecuación de una distribución simétrica a una no simétrica.

```
[22]: sns.boxplot(carsdf['price_usd'])
```

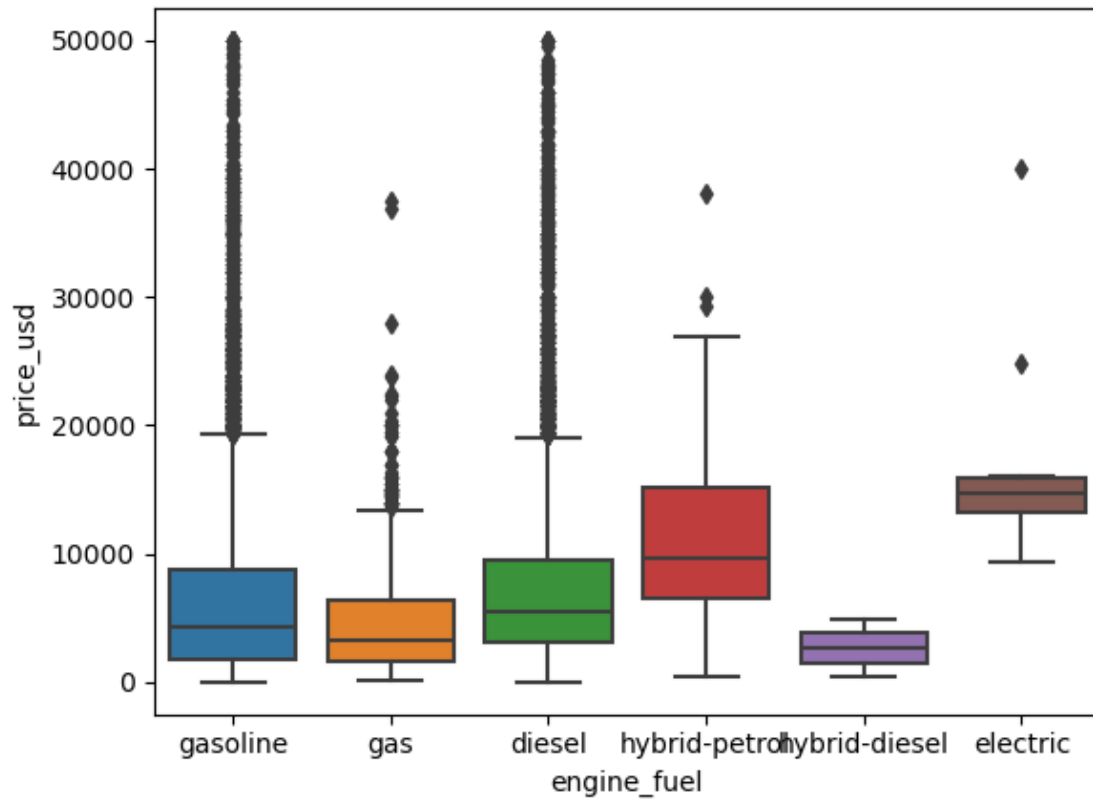
```
[22]: <Axes: >
```



```
[23]: sns.boxplot(x='engine_fuel', y='price_usd', data=carsdf)
```

```
[23]: <Axes: xlabel='engine_fuel', ylabel='price_usd'>
```





```
[24]: carsdf.dtypes
```

```
[24]: manufacturer_name    object
      model_name          object
      transmission        object
      color               object
      odometer_value      int64
      year_produced       int64
      engine_fuel         object
      engine_has_gas      bool
      engine_type         object
      engine_capacity     float64
      body_type           object
      has_warranty        bool
      state              object
      drivetrain          object
      price_usd           float64
      is_exchangeable     bool
      location_region     object
      number_of_photos    int64
      up_counter          int64
```

```

feature_0          bool
feature_1          bool
feature_2          bool
feature_3          bool
feature_4          bool
feature_5          bool
feature_6          bool
feature_7          bool
feature_8          bool
feature_9          bool
duration_listed    int64
dtype: object

```

```
[25]: carsdf.select_dtypes(include=['float64', 'int64'])
```

```

[25]:      odometer_value  year_produced  engine_capacity  price_usd  \
0           190000         2010           2.5      10900.00
1           290000         2002           3.0       5000.00
2           402000         2001           2.5       2800.00
3            10000         1999           3.0      9999.00
4           280000         2001           2.5       2134.11
...          ...          ...          ...          ...
38526        290000         2000           3.5       2750.00
38527        321000         2004           2.2       4800.00
38528        777957         2000           3.5       4300.00
38529         20000         2001           2.0       4000.00
38530        297729         2000           2.4       3200.00

```

```

      number_of_photos  up_counter  duration_listed
0                9         13         16
1               12         54         83
2                4         72        151
3                9         42         86
4               14          7          7
...          ...          ...          ...
38526             5         85         301
38527             4         20         317
38528             3         63         369
38529             7        156         490
38530             4         73         632

```

```
[38531 rows x 7 columns]
```

```

[26]: corr_matrix = carsdf.select_dtypes(include=['float64', 'int64']).
      ↪corr(method='pearson')

```

```
[27]: corr_matrix
```

```
[27]:
```

	odometer_value	year_produced	engine_capacity	price_usd	\
odometer_value	1.000000	-0.488679	0.105704	-0.421204	
year_produced	-0.488679	1.000000	0.005059	0.705511	
engine_capacity	0.105704	0.005059	1.000000	0.296597	
price_usd	-0.421204	0.705511	0.296597	1.000000	
number_of_photos	-0.143708	0.258180	0.106691	0.316859	
up_counter	-0.020961	0.007945	0.079152	0.057382	
duration_listed	-0.000428	-0.017001	0.080081	0.033524	


	number_of_photos	up_counter	duration_listed
odometer_value	-0.143708	-0.020961	-0.000428
year_produced	0.258180	0.007945	-0.017001
engine_capacity	0.106691	0.079152	0.080081
price_usd	0.316859	0.057382	0.033524
number_of_photos	1.000000	0.073891	-0.028255
up_counter	0.073891	1.000000	0.698116
duration_listed	-0.028255	0.698116	1.000000

```
[28]: import matplotlib.pyplot as plt
```

```
[29]: fig, ax = plt.subplots(nrows=1, ncols=1, figsize=(4, 4))

sns.heatmap(
    corr_matrix,
    annot      = True,
    cbar       = False,
    annot_kws  = {"size": 6},
    vmin       = -1,
    vmax       = 1,
    center     = 0,
    cmap       = sns.diverging_palette(20, 220, n=200),
    square     = True,
    ax         = ax
)
ax.set_xticklabels(
    ax.get_xticklabels(),
    rotation = 45,
    horizontalalignment = 'right',
)

ax.tick_params(labelsize = 8)
```



## 3 Tipos de escalamiento

### 3.1 Lineal

**Min-max:** hace una transformación para que los datos entren en el rango  $[-1, 1]$  o  $[0, 1]$  por medio de una fórmula. Es uno de los más usados. Funciona mejor para datos uniformemente distribuidos.

**Z-Score:** es uno de los más comunes porque está basado en el promedio y desviación estándar. Funciona mejor para datos distribuidos “normalmente” (forma de campana de Gauss).

```
[30]: X = carsdf['price_usd']
```

```
[31]: X
```

```
[31]: 0      10900.00
      1       5000.00
      2       2800.00
      3       9999.00
      4       2134.11
      ...
      38526    2750.00
      38527    4800.00
```

```

38528    4300.00
38529    4000.00
38530    3200.00
Name: price_usd, Length: 38531, dtype: float64

```

```
[32]: Z_X = (X - X.mean())/X.std()
```

```
[33]: Z_X
```

```

[33]: 0        0.662714
      1       -0.255123
      2       -0.597368
      3        0.522550
      4       -0.700957
      ...
38526    -0.605146
38527    -0.286236
38528    -0.364019
38529    -0.410689
38530    -0.535142
Name: price_usd, Length: 38531, dtype: float64

```

```
[34]: Z_X.mean()
```

```
[34]: -6.78621698787254e-17
```

```
[35]: Z_X.std()
```

```
[35]: 1.0000000000000016
```

```
[36]: MM_X = (X-X.min())/(X.max() - X.min())
```


```
[37]: sns.histplot(MM_X)
```

```

c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):

```

```
[37]: <Axes: xlabel='price_usd', ylabel='Count'>
```




```
[38]: MM2_X = (2*X-X.min()-X.max())/(X.max() - X.min())
```

```
[39]: sns.histplot(MM2_X)
```

```
c:\ProgramData\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
```

```
[39]: <Axes: xlabel='price_usd', ylabel='Count'>
```




### 3.2 Transformaciones no lineales en Python

```
[40]: X_es = X/10000
```

```
[41]: X_tan = np.tanh(X_es)
```

```
[42]: X_tan.hist()
```

```
[42]: <Axes: >
```



## 4 Procesamiento para variables categóricas

**Dummy:** es la representación más compacta que se puede tener de los datos. Es mejor usarla cuando los inputs son variables linealmente independientes (no tienen un grado de correlación significativo). Es decir, las cuando se sabe que las categorías son independientes entre sí. **One-hot:** es más general. Permite incluir categorías que no estaban en el dataset inicialmente. De forma que si se filtra una categoría que no estaba incluida, igual se pueda representar numéricamente.

```
[43]: ZZ = pd.get_dummies(carsdf['engine_type'])
```

```
[44]: ZZ
```

```
[44]:
```

	diesel	electric	gasoline
0	False	False	True
1	False	False	True
2	False	False	True
3	False	False	True
4	False	False	True
...	...	...	...
38526	False	False	True
38527	True	False	False
38528	False	False	True



```
38529    False    False    True
38530    False    False    True
```

```
[38531 rows x 3 columns]
```

```
[45]: import sklearn.preprocessing as preprocessing
```

```
encoder = preprocessing.OneHotEncoder(handle_unknown='ignore')
```

```
[46]: tmp = carsdf['engine_type'].unique()
```

```
[47]: tmp
```

```
[47]: array(['gasoline', 'diesel', 'electric'], dtype=object)
```

```
[48]: encoder.fit(carsdf[['engine_type']].values)
```

```
[48]: OneHotEncoder(handle_unknown='ignore')
```

```
[49]: encoder
```

```
[49]: OneHotEncoder(handle_unknown='ignore')
```

```
[50]: encoder.transform(['gasoline'], ['diesel'], ['aceite'], ['agua']).toarray()
```

```
[50]: array([[0., 0., 1.],
           [1., 0., 0.],
           [0., 0., 0.],
           [0., 0., 0.]])
```

```
[51]: encoder.transform(['gasoline'], ['diesel'], ['aceite']).toarray()
```

```
[51]: array([[0., 0., 1.],
           [1., 0., 0.],
           [0., 0., 0.]])
```

## 5 Álgebra Lineal aplicada a la matriz de correlaciones

```
[52]: corr_matrix
```

```
[52]:
```

	odometer_value	year_produced	engine_capacity	price_usd	\
odometer_value	1.000000	-0.488679	0.105704	-0.421204	
year_produced	-0.488679	1.000000	0.005059	0.705511	
engine_capacity	0.105704	0.005059	1.000000	0.296597	
price_usd	-0.421204	0.705511	0.296597	1.000000	
number_of_photos	-0.143708	0.258180	0.106691	0.316859	

up_counter	-0.020961	0.007945	0.079152	0.057382
duration_listed	-0.000428	-0.017001	0.080081	0.033524

	number_of_photos	up_counter	duration_listed
odometer_value	-0.143708	-0.020961	-0.000428
year_produced	0.258180	0.007945	-0.017001
engine_capacity	0.106691	0.079152	0.080081
price_usd	0.316859	0.057382	0.033524
number_of_photos	1.000000	0.073891	-0.028255
up_counter	0.073891	1.000000	0.698116
duration_listed	-0.028255	0.698116	1.000000

```
[53]: from numpy.linalg import eig
```

```
[54]: w,v=eig(corr_matrix)
```

```
[55]: w
```

```
[55]: array([2.26276437, 1.71228239, 1.12619392, 0.83873606, 0.53102881,
          0.23486557, 0.29412887])
```

```
[56]: w.sum()
```

```
[56]: 7.0000000000000007
```

```
[57]: w # valores propios
```

```
[57]: array([2.26276437, 1.71228239, 1.12619392, 0.83873606, 0.53102881,
          0.23486557, 0.29412887])
```

```
[58]: w
```

```
[58]: array([2.26276437, 1.71228239, 1.12619392, 0.83873606, 0.53102881,
          0.23486557, 0.29412887])
```

```
[59]: porc_varianza = w/w.sum()
```

```
[60]: porc_varianza
```

```
[60]: array([0.32325205, 0.24461177, 0.16088485, 0.11981944, 0.07586126,
          0.03355222, 0.04201841])
```

```
[61]: porc_varianza
```

```
[61]: array([0.32325205, 0.24461177, 0.16088485, 0.11981944, 0.07586126,
          0.03355222, 0.04201841])
```

```
[62]: porc_varianza.sum()
```

```
[62]: 0.9999999999999999
```

```
[63]: V_pd = pd.DataFrame.from_records(v)
```

```
[64]: V_pd.T.dot(V_pd)
```

```
[64]:
```

	0	1	2	3	4	\
0	1.000000e+00	-1.352229e-16	-3.328643e-16	8.855268e-17	6.918030e-17	
1	-1.352229e-16	1.000000e+00	4.073029e-16	-6.527008e-17	-6.617097e-16	
2	-3.328643e-16	4.073029e-16	1.000000e+00	1.627564e-16	1.207682e-16	
3	8.855268e-17	-6.527008e-17	1.627564e-16	1.000000e+00	-7.056217e-16	
4	6.918030e-17	-6.617097e-16	1.207682e-16	-7.056217e-16	1.000000e+00	
5	9.916572e-17	1.877511e-16	-8.315955e-17	-1.371748e-16	-3.246695e-16	
6	5.706270e-17	-4.578607e-16	-1.785732e-16	5.362632e-16	8.951776e-17	

	5	6
0	9.916572e-17	5.706270e-17
1	1.877511e-16	-4.578607e-16
2	-8.315955e-17	-1.785732e-16
3	-1.371748e-16	5.362632e-16
4	-3.246695e-16	8.951776e-17
5	1.000000e+00	9.615936e-16
6	9.615936e-16	1.000000e+00